













**Your submission is being evaluated - please don't share these result links publicly (e.g. on Ed Discussion).**







Test setup: All VM's pinned to the same CPU. Equal load on all vCPUs







Test setup: 50% of VM's pinned to pCPU 0 and the rest to pCPU 1. Equal load on all vCPUs







Test setup: Balanced pinning of vCPU to pCPU. Equal load on all vCPUs







Test setup: Each vCPU has equal affinity to all pCPUs. Equal load on all vCPUs



Test setup: Each vCPU has equal affinity to all pCPUs. 50% vCPUs run heavy load and the remaining run a light load







Test setup: Balanced pinning of vCPU to pCPU. No load on any vCPUs







Test setup: Each vCPU has equal affinity to all pCPUs. All vCPUs run light load, 50% switch to heavy load later





**Your submission is being evaluated - please don't share these result links publicly (e.g. on Ed Discussion).**







Test Setup: 1st VM requests memory until limit followed by releasing all memory.



Test Setup: All VMs request memory until limit followed by them releasing the requested memory.



Test Setup: All VMs begin requesting memory. First 2 VMs start releasing memory while the other 2 keep consuming.



Test Setup: First 2 VMs begin consuming memory. The first VM starts releasing memory, while the second continues consuming.